

ENTERPRISE AI KNOWLEDGE ASSISTANT

Enterprise Project Management Toolkit

Executive Brief

Press Release & Frequently Asked Questions (PR-FAQ)

Document Title	Executive Brief & PR-FAQ
Project Name	Enterprise AI Knowledge Assistant (EAIKA)
Document Owner	Mautik Patel, AI Project Delivery Manager
Version	1.0
Sponsor	John Whitfield, Chief AI Officer
Status	Approved for Distribution
Classification	Internal – Executive Stakeholders
Date	July 4, 2026

1. Purpose of this Document

This Executive Brief introduces the Enterprise AI Knowledge Assistant (EAIKA) to sponsors, steering committee members, and business stakeholders using the “working backwards” PR-FAQ format. It is written as if the product has already launched, describing the customer problem, the solution, and the value delivered, followed by anticipated questions from executives, security teams, and end users.

2. Press Release

Enterprise AI Knowledge Assistant Turns Scattered Company Documents into Instant, Trustworthy Answers

Organizations today manage thousands of documents across **project management, HR, finance, operations, and compliance**. Employees often spend hours every week searching through folders, emails, and shared drives or reaching out to colleagues for information that already exists somewhere within the organization. The **Enterprise AI Knowledge Assistant** addresses this challenge by allowing users to ask questions in **plain language** and receive **accurate, source-backed answers** within seconds.

The platform is a **locally deployed Retrieval-Augmented Generation (RAG) solution** built using **LangChain, Ollama, ChromaDB, and Streamlit**, providing a secure and intuitive web-based experience. Documents in **PDF, Word, PowerPoint, Excel, text, CSV, email, and image** formats are uploaded once, automatically processed into a **searchable knowledge base**, and made available for **natural language question answering**. Every response includes **references to the original source document**, allowing users to verify the information with confidence.

Unlike cloud-based AI assistants, the **Enterprise AI Knowledge Assistant** operates entirely within the organization's own environment. **No document content or user queries are transmitted to external AI providers**, making the platform well suited for **confidential project documentation, HR records, compliance materials, and other sensitive enterprise information**.

"Knowledge already exists inside the organization. The real challenge has always been finding the right information at the right time. This platform solves that problem without requiring organizations to compromise on data privacy." - *Product Sponsor*

The initial release focuses on **Project and Program Management** as the primary use case. Delivery teams, project managers, and executives can ask questions such as **"What is our current project health?"** or **"What are the highest project risks?"** and receive **immediate, cited responses** generated from **project charters, status reports, RAID logs, UAT records, and other project documentation**. Future releases will expand the platform to support **HR, Finance, Compliance, Healthcare, and Operations**, creating a **unified enterprise knowledge platform** across multiple business domains.

3. Customer / Internal FAQ

Q1: What problem does this solve?

Enterprise knowledge is scattered across hundreds of documents, formats, and repositories. Finding an answer today requires manually opening files, searching email threads, or relying on a subject matter expert who may not be available. This leads to lost productivity, inconsistent answers, and delayed decisions.

Q2: How is this different from a simple keyword search tool (e.g., SharePoint search)?

The assistant uses semantic search rather than keyword matching. It understands the meaning behind a question, so a user asking, **"Is the system healthy?"** will correctly retrieve a document that instead states, **"application availability remains within acceptable service levels,"** even though no words overlap. Keyword search cannot make that connection.

Q3: Where does our data go? Is this secure?

All processing, including document ingestion, embedding generation, vector storage, and language model inference, takes place entirely within the organization's local environment. The platform uses Ollama-hosted open-weight models (Qwen2.5) together with a local ChromaDB vector store to process and retrieve information securely. No document content or user queries are transmitted to third-party cloud AI services, significantly reducing data residency, privacy, and confidentiality risks when compared to cloud-hosted AI assistants.

Q4: What document types are supported at launch?

PDF, DOCX, PPTX, XLSX, TXT, CSV, EML (email), and image formats (PNG, JPG, JPEG) with OCR text extraction for scanned or image-based content.

Q5: Who is the first internal audience for this tool?

Project and program management teams are the flagship audience: PMO leads, project managers, delivery leads, and executive sponsors who need instant, source-backed answers about project health, risk, scope, and status without manually searching charters, status reports, and RAID logs.

Q6: What does a typical user interaction look like?

A user opens the Streamlit web application, types a natural-language question such as “What are the highest project risks?” and receives a concise answer along with references to the specific source document(s) and passages the answer was generated from, allowing the user to verify accuracy before acting on it.

Q7: What administrative controls exist?

An Administrator Dashboard allows authorized users to upload new documents, trigger knowledge base builds, rebuild the vector database when documents change, and review knowledge base statistics and document inventory.

Q8: What is out of scope for the initial release?

Real-time integration with live enterprise systems (e.g., direct ERP or ticketing system connections), automated write-back actions, and multi-tenant access controls are out of scope for this release and are captured as future roadmap items in the Product Requirements Document.

Q9: What is the anticipated business value?

Reduced time spent searching for information, improved consistency of answers across teams, faster onboarding of new employees, preserved institutional knowledge, and higher confidence in AI adoption due to the fully local, privacy-first architecture.

4. Executive Summary Table

Dimension	Summary
Problem	Enterprise knowledge is fragmented across formats and repositories, costing employee time, and creating inconsistent answers.
Solution	A locally hosted RAG platform (LangChain + Ollama + ChromaDB + Streamlit) providing natural-language Q&A over uploaded enterprise documents, with cited sources.
Primary Launch Use Case	Project & Program Management knowledge search (charters, BRDs, status reports, RAID logs, UAT signoffs).
Differentiator	100% local deployment — no document or query data leaves the organization's environment.
Target Users	PMO, project managers, delivery leads, business stakeholders, and (future phases) HR, Finance, Compliance, and Operations teams.
Success Measure	Reduction in time-to-answer for enterprise knowledge questions and adoption across at least one pilot department within the first quarter.

5. Approval

This Executive Brief is submitted for sponsor review and approval prior to formal Project Charter development and stakeholder kickoff.