

ENTERPRISE AI KNOWLEDGE ASSISTANT

Enterprise Project Management Toolkit

Project Charter

Authorization to Plan and Execute the EAIKA Program

Document Title	Project Charter
Project Name	Enterprise AI Knowledge Assistant (EAIKA)
Document Owner / Project Manager	Mautik Patel, AI Project Delivery Manager
Version	1.0
Project Sponsor	John Whitfield, Chief AI Officer
Status	Approved
Date	July 4, 2026

1. Project Purpose

This charter formally authorizes the Enterprise AI Knowledge Assistant (EAIKA) project and grants the Project Manager the authority to apply organizational resources to project activities. EAIKA is a locally deployed, Retrieval-Augmented Generation (RAG) platform that enables employees to search and query enterprise documents using natural language and receive accurate, source-cited answers.

2. Business Case

Enterprise knowledge is fragmented across project documentation, policies, reports, presentations, and email, forcing employees to manually search multiple systems to find information that already exists. This results in lost productivity, inconsistent answers across teams, delayed decision-making, and heavy reliance on individual subject matter experts. EAIKA addresses this by consolidating enterprise knowledge into a single, AI-searchable platform while keeping all data processing on-premises to satisfy confidentiality requirements.

3. Project Objectives

- Deliver a working RAG-based knowledge assistant covering the Project & Program Management document domain as the flagship use case.
- Enable natural-language, source-cited question answering across PDF, DOCX, PPTX, XLSX, TXT, CSV, EML, and image (OCR) document formats.
- Ensure 100% local processing with no document content transmitted to third-party AI services.
- Provide an administrator dashboard for document upload, knowledge base build/rebuild, and inventory management.
- Establish a reusable architecture that can be extended to HR, Finance & Compliance, Healthcare, and Operations knowledge domains in future phases.

4. Scope Summary

In Scope

- Document ingestion pipeline (upload, chunking, embedding, indexing) for the supported formats listed above.
- Semantic search and retrieval using ChromaDB as the vector store.
- Local LLM inference via Ollama (Qwen2.5) for answer generation, with source attribution.
- Streamlit-based end-user chat interface and administrator dashboard.
- Pilot deployment for a project management document set (charters, BRDs, status reports, RAID logs, UAT records).

Out of Scope (Phase 1)

- Real-time bidirectional integration with live enterprise systems (ERP, ITSM, ticketing).
- Multi-tenant role-based access control and per-document per missioning.
- Cloud-hosted or multi-region deployment.
- Automated workflow actions (e.g., auto updating a ticket) beyond information retrieval.

5. High-Level Milestones

Milestone	Target Timeframe	Owner	Description
Project Charter & Funding Approval	Jan-26	CRO / Sponsor	Project Charter approved, funding released, and project officially initiated.
Requirements Sign-off (BRD, FRD & PRD)	Feb-26	Mautik Patel	Business, functional, and product requirements reviewed and approved by the sponsor and key stakeholders.
Architecture & Technical Design Approval	Mar-26	Architecture Review Board	Technical Design Document reviewed and approved, confirming the solution architecture and technology stack.

Core Platform Build Complete	Apr-26	Data Engineering	Document ingestion pipeline, embedding generation, vector database, retrieval pipeline, and AI services developed and integrated.
UI & Administrator Dashboard Complete	May-26	Engineering Team	Streamlit-based chat interface and administrator dashboard completed and ready for system testing.
Model Validation & AI Safety Sign-off	Jun-26	Data Science & Compliance	AI model validated for performance, security, compliance, and responsible AI requirements.
User Acceptance Testing (UAT) Sign-off	Jul-26	Business Stakeholders	User Acceptance Testing completed successfully with formal business approval.
Pilot Deployment	Aug-26	Project Team	Solution deployed to the pilot department for controlled production validation and user feedback.
Production Go-Live	Sep-26	Mautik Patel	Enterprise production deployment completed, and solution made available to end users.
Post-Implementation Review	Dec-26	Mautik Patel	Project outcomes reviewed, lessons learned documented, and success metrics evaluated after production rollout.

6. Budget Summary

The project is resourced as an internal engineering initiative leveraging open-source components (LangChain, ChromaDB, Ollama, Streamlit) to minimize licensing cost. Primary costs are internal labor (engineering, QA, PMO), on-premises or private-cloud compute for LLM hosting, and infrastructure for the vector database and application server. Detailed cost estimates are maintained in the Infrastructure & Capacity Plan.

Category	Amount (USD)	% of Total
Data science & model development (staff)	\$820,000	34%
Engineering & integration (staff)	\$610,000	25%
Cloud/Snowflake compute & storage	\$340,000	14%
Third-party fraud data & model tooling licenses	\$280,000	12%
Model risk validation (external)	\$150,000	6%
Contingency (9%)	\$200,000	9%
Total	\$2,400,000	100%

7. Key Stakeholders

Stakeholder	Role	Department / Organization	Key Responsibilities	Influence
Sarah Whitfield	Executive Sponsor	Executive Leadership	Provides executive sponsorship, approves project funding, removes organizational roadblocks, and ensures alignment with business strategy.	High
Mautik Patel	Project Delivery Manager	PMO / Enterprise Delivery	Leads project planning, execution, budget management, stakeholder communication, governance, risk management, and successful project delivery.	High
Mautik Patel	Product Owner	Business Operations	Defines business requirements, prioritizes the product backlog, validates delivered features, and accepts completed functionality.	High
Michael Chen	Engineering Lead	Engineering	Leads solution architecture, document ingestion, vector database implementation, AI integration, and overall technical delivery.	High

Priya Nataraj	Head of Data Science	Data & Analytics	Oversees AI model development, embedding strategy, evaluation, performance optimization, and model validation.	High
Emma Rodriguez	QA & Test Lead	Quality Assurance	Develops the testing strategy, manages test execution, coordinates User Acceptance Testing (UAT), and ensures product quality.	Medium
David Brooks	Information Security Officer	Information Security	Reviews security architecture, access controls, data privacy, local deployment, and cybersecurity compliance.	High
Olivia Carter	Legal & Compliance Manager	Legal & Compliance	Ensures regulatory compliance, AI governance, risk management, and adherence to organizational policies.	High
Daniel Murphy	Infrastructure Manager	Enterprise IT	Provides infrastructure, networking, storage, production environment support, monitoring, backup, and operational readiness.	Medium
Rachel Kim	Business Representative	Project & Program Management Office	Represents end users, participates in User Acceptance Testing (UAT), validates business scenarios, and provides feedback before production rollout.	Medium

8. Assumptions

- Pilot department documents (project management artifacts) will be made available in supported formats for ingestion.
- Adequate on-premises or private compute capacity exists or can be provisioned to host the local LLM and vector database.
- Business stakeholders will be available for requirements validation and UAT within the planned schedule.

9. Constraints

- All AI inference and document processing must remain within the organization's-controlled environment; no external cloud AI APIs may be used for document content.
- The initial release is limited to the document formats and use case defined in the Scope Summary.
- Delivery timeline is fixed at 13 weeks to production go-live, with a 4-week post-launch stabilization and review window.

10. High-Level Risks

Detailed risks are tracked in the RAID Log & Risk Register and the AI Risk Evaluation & Safety Case. Top charter-level risks include potential for the LLM to generate an inaccurate or unsupported answer ("hallucination"), insufficient compute capacity for acceptable response latency, and slower-than-expected user adoption if trust in AI-generated answers is not established early through strong source citation.

11. Approval

Name / Role	Signature	Date
Sarah Whitfield – Chief Risk Officer (Sponsor)		Jan 12, 2026
Mautik Patel – Project Delivery Manager		Jan 12, 2026
Priya Nataraj – Head of Data Science		Jan 12, 2026