

ENTERPRISE AI KNOWLEDGE ASSISTANT

Enterprise Project Management Toolkit

Product Requirements Document

PRD – Enterprise AI Knowledge Assistant

Document Title	Product Requirements Document (PRD)
Project Name	Enterprise AI Knowledge Assistant (EAIKA)
Product Owner	AI Platform Product Team
Version	1.0
Status	Approved
Date	July 4, 2026

1. Product Vision

Enterprise AI Knowledge Assistant (EAIKA) is a **private, fully local Retrieval-Augmented Generation (RAG) platform** that enables organizations to transform enterprise documents into an intelligent, conversational knowledge base. Users can ask questions in **natural language** and receive **concise, accurate, source-cited answers**, without any document content leaving the organization's own environment.

2. Product Overview

Attribute	Detail
Product Type	Local RAG-based enterprise knowledge assistant
Delivery Interface	Streamlit web application
Core AI Stack	LangChain (orchestration), Ollama (local LLM runtime), Qwen2.5 (language model), ChromaDB (vector database)
Deployment Model	100% on-premises / private environment; no external cloud AI dependency
Primary Launch Domain	Project & Program Management knowledge (charters, BRDs, status reports, RAID logs, UAT records)
Future Domains	HR & Company Policies, Healthcare & Clinical Operations, Banking/Finance & Compliance, Manufacturing & Operations, Legal & Contracts, Training & Internal Knowledge

3. Target Personas

Priya – Project Manager

Needs a fast way to answer stakeholder questions about project status, risks, and milestones without opening five different documents before a steering committee meeting.

Marcus – Executive Sponsor

Wants a quick, trustworthy summary of program health across multiple projects without waiting for a formal report to be compiled.

Elena – Knowledge Base Administrator

Responsible for keeping the knowledge base current: uploading new documents as they are produced and monitoring indexing health.

4. Feature Set

4.1 AI Enterprise Search

- Ask questions in natural language instead of manually searching documents.
- Semantic, context-aware answers powered by the local Qwen2.5 model.
- Every answer accompanied by source-backed references.

4.2 Intelligent Knowledge Base

- Document upload with automatic chunking and embedding generation.
- ChromaDB-based indexing for fast semantic retrieval.

4.3 Semantic Search

- Retrieves information based on meaning, not just exact keyword overlap, so varying terminology across teams does not block discovery.

4.4 Multi-Format Support

- PDF, DOCX, PPTX, XLSX, TXT, CSV, EML, PNG, JPG, JPEG, with OCR for image-based content.

4.5 Source Transparency

- Users can verify information, review original content, and navigate back to the source document — avoiding a “black box” AI experience.

4.6 Administrator Dashboard

- Upload documents, build or rebuild the knowledge base, review document statistics and inventory.

4.7 Privacy-First Architecture

- Runs completely offline with local AI models; no cloud APIs; no data leaves the organization's machine or private network.

5. User Stories & Acceptance Criteria

User Story	Acceptance Criteria
As a project manager, I want to ask, “What is the current project health?” so that I can quickly prepare for a status meeting.	Given relevant status reports are indexed, when the user submits the question, then the system returns a concise health summary with citations to the source status report(s) within an acceptable response time.
As an executive sponsor, I want to ask, “What are the highest project risks?” so that I can prioritize mitigation discussions.	The system retrieves risk-related content from the RAID Log / Risk Register and returns a ranked summary with source citations.
As a knowledge base administrator, I want to upload a new batch of documents so that the assistant stays current.	Given valid supported files, when uploaded via the Administrator Dashboard, then the documents appear in the inventory and are available for query after a successful build.
As any user, I want to see which document an answer came from so that I can trust and verify the response.	Every generated answer displays at least one source document reference; the user can access the referenced content.
As an administrator, I want to see knowledge base statistics so that I can monitor system health.	The dashboard displays total documents, total chunks, and last build/rebuild time, refreshed after each build operation.

6. Non-Goals (Phase 1)

- Automated write-back to source systems (e.g., updating a project plan directly).
- Fine-grained, per-document role-based access control.
- Multi-tenant SaaS deployment across separate customer organizations.
- Real-time streaming ingestion from live systems (e.g., ticketing systems, email inboxes) - ingestion is upload-based in Phase 1.

7. Success Metrics

Metric	Target
Answer accuracy (UAT sample)	≥ 80% of test questions answered correctly with valid source citation
Average response latency	Within agreed enterprise readiness SLA thresholds for the deployed hardware profile
Pilot department adoption	Active use by pilot department within 30 days of go-live
User-reported trust score	Positive trend in post-launch satisfaction survey
Zero external data egress	No document content transmitted to external/cloud AI services, confirmed via security review

8. Release Plan

Release	Scope
Phase 1 (this PRD)	Project & Program Management document domain; core RAG pipeline; Streamlit chat UI; Administrator Dashboard.
Phase 2 (future)	HR & Company Policies domain; expanded document connectors; usage analytics.
Phase 3 (future)	Finance & Compliance and Healthcare domains; role-based access control; audit logging enhancements.

9. Open Questions

- What hardware profile (CPU/GPU) will host the local LLM in production, and does it meet target latency for the pilot user base?
- Will document versioning be required so outdated documents can be retired from the knowledge base without manual cleanup?
- What is the process for handling documents that fail OCR or text extraction?

10. Approval

This PRD requires sign-off from the Product Owner, Engineering Lead, and Business Sponsor before entering the design phase.